

TensorFlow

- MNIST example -

2019 – 2020

Ando Ki, Ph.D.

adki@future-ds.com

Table of contents

- Get TensorFlow packages
- MNIST using MLP
- Project: MNIST using one hidden layer
- MNIST using CNN

Get TensorFlow packages

- This package is required to utilize MNIST dataset.

- Go to the project directory

```
$ cd $(PROJECT)/codes
```

- Get a copy of TensorFlow package

```
$ git clone https://github.com/tensorflow/tensorflow.git
```

- ▶ or visit following site and get a copy of it
 - ➔ check directory hierarchy and its name; modify if necessary


```
https://github.com/tensorflow/tensorflow
```

3

MNIST using MLP

- MLP (Multi-Layer Perceptron) applies to MNIST

- ▶ 784 (28x28) inputs of black and white → converted to floating number 0.0 ~ 1.0
- ▶ 10 outputs representing digit 0 to 9
- ▶ Two hidden layer
- ▶ 256 features for each hidden layer



```
$(PROJECT)/codes/tensorflow-projects/mnist-projects/mlp
```

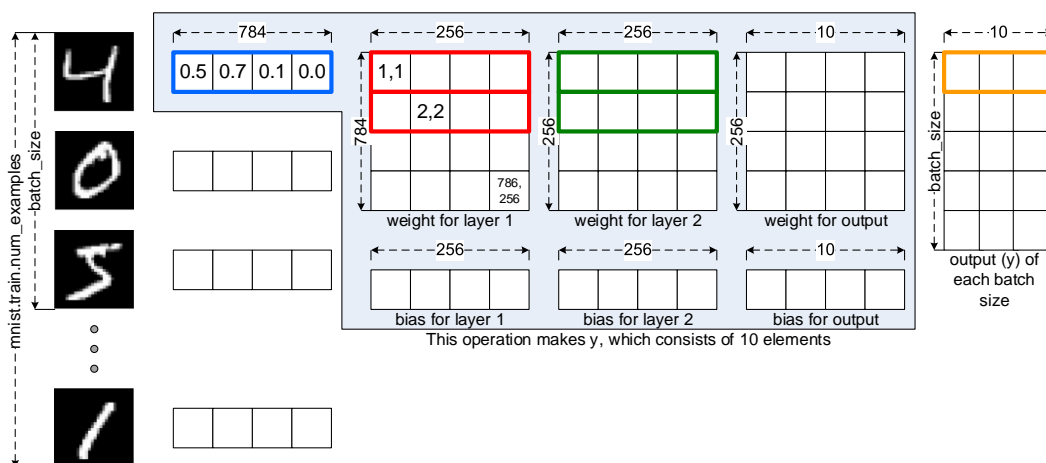
4

MNIST using MLP

- Total testing patterns
 - ▶ 60K images (train patterns: 55K images)
- Batch number
 - ▶ 100 → group 600 batches (train batches: 550)
- Epoch number
 - ▶ The number of whole trainings (forward calculation and back propagation)

5


MNIST using MLP



6

Terminology

- Iteration
 - ▶ Forward and backward for a number of inputs (i.e., batch or minibatch)
 - Batch or minibatch
 - ▶ A number of inputs (e.g., testing images) to complete an iteration
 - ▶ Batch size: the number of training examples in one forward/backward pass
 - ☞ The bigger the batch size, the more memory space needed
 - ▶ Minibatch
 - ☞ Take a small number of examples at a time, ranging from 1 to a few hundred, during one iteration
 - Epoch
 - ▶ one epoch: one forward pass and one backward pass of all the training examples
 - ▶ it contains a number of iterations
- 1 epoch = (number of iterations)
 - = (total training examples) /
 - (batch size or minibatch size)



7

TensorFlow operators

- `tf.placeholder()`
- `tf.Variable()`
- `tf.global_variables_initializer()`
- `tf.add()`
- `tf.matmul()`
- `tf.reduce_mean()`
- `tf.nn.relu()`
- `tf.train.AdamOptimizer()`
- `tf.train.GradientDescentOptimizer()`
- `tf.argmax()`

8

MNIST using MLP

\$(PROJECT)/codes/tensorflow-projects/mnist-projects/mlp/mnist_mlp.py

```
#-----
import os
import sys
os.environ['TF_CPP_MIN_LOG_LEVEL']='2'
sys.path.append(os.path.dirname("../..../tensorflow/tensorflow"))
#-----
import tensorflow as tf
import tensorflow as tf
import numpy as np
from random import randint
import matplotlib.pyplot as plt
#-----
# Import MNIST data
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("../dataset", one_hot=True)
#-----
# Parameters
learning_rate = 0.01
training_epochs = 1000
batch_size = 100
display_step = 10
```

Depress warnings

Path where to find module

TensorFlow module

MNIST dataset handling module

Read MNIST dataset

Parameters



9

MNIST using MLP

```
#-----
# tf Graph Input
x = tf.placeholder("float", [None, 784], name="x-input")
y = tf.placeholder("float", [None, 10], name="y-output")
#-----
# Model weights
W1 = tf.Variable(tf.random_normal([784,256]), name="weight1")
b1 = tf.Variable(tf.random_normal([256]), name="bias1")
W2 = tf.Variable(tf.random_normal([256,256]), name="weight2")
b2 = tf.Variable(tf.random_normal([256]), name="bias2")
W3 = tf.Variable(tf.random_normal([256,10]), name="weight-out")
b3 = tf.Variable(tf.random_normal([10]), name="bias-out")
#-----
# inference -> hypothesis
layer_1 = tf.add(tf.matmul(x, W1), b1)
layer_1 = tf.nn.sigmoid(layer_1, name="layer1-sigmoid")
layer_2 = tf.add(tf.matmul(layer_1, W2), b2)
layer_2 = tf.nn.sigmoid(layer_2, name="layer2-sigmoid")
out_layer = tf.add(tf.matmul(layer_2, W3), b3)
infer = tf.nn.softmax(out_layer, name="infer")
```

Input and output

Network parameters



Network building

Note softmax

10

MNIST using MLP

```
#-----
# Minimize error using cross entropy
cost = tf.reduce_mean(-tf.reduce_sum(y * tf.log(infer), reduction_indices=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)

#-----
# Evaluate -- 0~1 accuracy
correct = tf.equal(tf.argmax(infer, 1), tf.argmax(y,1))
accuracy = tf.reduce_mean(tf.cast(correct, "float"))

#-----
# Initializing the variables
init = tf.global_variables_initializer()

#-----
# Add ops to save and restore all the variables
saver = tf.train.Saver()
```

cost and optimizer

y => (100, 10)
cost will be scalar

Accuracy

Initialize

Prepare for save/restore

11

MNIST using MLP

```
with tf.Session() as sess:
    sess.run(init)
    # tensorboard --logdir=./logs
    # http://localhost:6006
    write = tf.summary.FileWriter('./logs', sess.graph)
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.
        total_batch = int( mnist.train.num_examples / batch_size )
        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            # Fit training using batch data
            sess.run(optimizer, feed_dict={x: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += sess.run(cost, feed_dict={x: batch_xs, y: batch_ys})
            avg_accu += sess.run(accuracy, feed_dict={x: batch_xs, y: batch_ys})

        avg_cost /= total_batch
        avg_accu /= total_batch
        # Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", "%04d" % (epoch + 1), "cost=", "{:.9f}".format(avg_cost),\
                  "accuracy=", "{:.6f}".format(avg_accu)

        if avg_accu > 0.95:
            break
```

training epochs = 100

total_batch = 550

batch_xs = (100, 784)
batch_ys = (100, 10)

Stop when sufficient accuracy has been reached

12

MNIST using MLP

```
with tf.Session() as sess:
    sess.run(init)
    # tensorboard --logdir=./logs
    # http://localhost:6006
    write = tf.summary.FileWriter('./logs', sess.graph)
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.
        total_batch = int( mnist.train.num_examples / batch_size )
        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            # Fit training using batch data
            sess.run(optimizer, feed_dict={x: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += sess.run(cost, feed_dict={x: batch_xs, y: batch_ys})
            avg_accu += sess.run(accuracy, feed_dict={x: batch_xs, y: batch_ys})

        avg_cost /= total_batch
        avg_accu /= total_batch
        # Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", "%04d" % (epoch + 1), "cost=", "{:.9f}".format(avg_cost),\
                  "accuracy=", "{:.6f}".format(avg_accu)

        if avg_accu>0.95:
            break
```

epoch times over all
data set

for each batch

13

MNIST using MLP

```
# Save weights
saver.save(sess, "./model/model.ckpt")

#-----
# Test model
print "Accuracy:", accuracy.eval({x: mnist.test.images, y: mnist.test.labels})

#-----
# predict & show
r = randint(0, mnist.test.num_examples - 1)
print "Label: ", sess.run(tf.argmax(mnist.test.labels[r:r+1], 1))
print "Prediction: ", sess.run(tf.argmax(infer, 1), {x: mnist.test.images[r:r+1]})
# show the img
plt.imshow(mnist.test.images[r:r+1].reshape(28, 28),
           cmap="Greys", interpolation="nearest")
plt.show()
#-----
```

Save network parameters

Check accuracy for validation set

Pick a number between 0 ~ num_examples

Print which label

Inference


Show image

14

MNIST using MLP

```
(tensorflow)$ python mnist_mlp.py
Extracting ../dataset/train-images-idx3-ubyte.gz
Extracting ../dataset/train-labels-idx1-ubyte.gz
Extracting ../dataset/t10k-images-idx3-ubyte.gz
Extracting ../dataset/t10k-labels-idx1-ubyte.gz
Epoch: 0001 cost= 5.25198 accuracy= 0.30315
Epoch: 0011 cost= 0.89775 accuracy= 0.77602
Epoch: 0021 cost= 0.64263 accuracy= 0.82938
Epoch: 0031 cost= 0.52474 accuracy= 0.85504
Epoch: 0041 cost= 0.45274 accuracy= 0.87211
Epoch: 0051 cost= 0.40308 accuracy= 0.88496
Epoch: 0061 cost= 0.36560 accuracy= 0.89400
.....
Epoch: 0161 cost= 0.20485 accuracy= 0.93944
Epoch: 0171 cost= 0.19650 accuracy= 0.94225
Epoch: 0181 cost= 0.18897 accuracy= 0.94424
Epoch: 0191 cost= 0.18194 accuracy= 0.94656
Epoch: 0201 cost= 0.17524 accuracy= 0.94880
Optimization finished
Accuracy: 0.9113
Label: [5]
Prediction: [5]
^C
(tensorflow)$ tensorboard --logdir=logs
```

When 'matplotlib' is missing:
 \$ sudo apt-get install libpng-dev
 \$ sudo apt-get install libfreetype6-dev
 \$ pip install matplotlib



Training accuracy

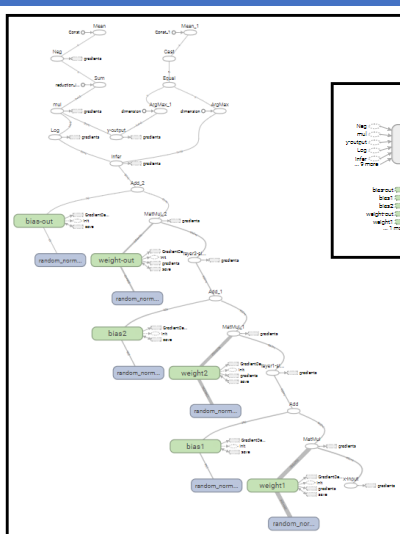
Validation accuracy

Predict correct label


Invoke tensor board

15

MNIST using MLP



invoke web browser: <http://localhost:6006>
 Then select 'GRAPH' menu



16

MNIST using MLP inference

\$(PROJECT)/codes/tensorflow-projects/mnist-projects/mlp/mnist_inference.py

```
#-----
import os
import sys
os.environ['TF_CPP_MIN_LOG_LEVEL']='2'
sys.path.append(os.path.dirname("../..../tensorflow/tensorflow"))

#-----
import tensorflow as tf
from random import randint
import matplotlib.pyplot as plt

#-----
sess = tf.Session()

#-----
# Create the network
saver = tf.train.import_meta_graph("../model/model.ckpt.meta")

#-----
# Restore parameters
saver.restore(sess, "../model/model.ckpt")
```

Build network from save data file

Restore network parameters from saved data file

17

MNIST using MLP inference

```
#-----
# get references of graph and tensors
graph = tf.get_default_graph()
x = graph.get_tensor_by_name("x-input:0");
y = graph.get_tensor_by_name("y-output:0");
infer = graph.get_tensor_by_name("infer:0")

#-----
# testing data-set
#-----
# Import MNIST data
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("../dataset", one_hot=True)
```

Get graph

Get input and output

Get inference operator node

MNIST dataset handling module

Read MNIST dataset

18

MNIST using MLP inference

```
#-----
# predict & show
for i in range(10):
    r = randint(0, mnist.test.num_examples - 1)
    print "Label: ", sess.run(tf.argmax(mnist.test.labels[r:r+1], 1))
    print "Prediction: ", sess.run(tf.argmax(infer, 1),\
                                   {x: mnist.test.images[r:r+1]})

#-----
# show the img
plt.imshow(mnist.test.images[r:r+1].reshape(28, 28),\
           cmap="Greys", interpolation="nearest")
plt.show()

#-----
```

Test 10 patterns

Pick a number between 0 ~ num_examples

Print which label

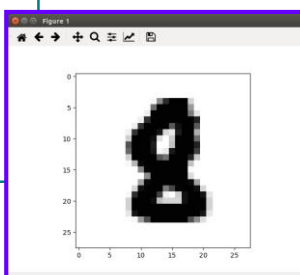
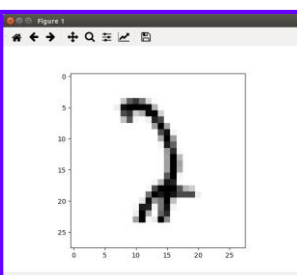

Inference

Show image

19

MNIST using MLP inference

```
(tensorflow)$ python mnist_inference.py
Extracting ../dataset/train-images-idx3-ubyte.gz
Extracting ../dataset/train-labels-idx1-ubyte.gz
Extracting ../dataset/t10k-images-idx3-ubyte.gz
Extracting ../dataset/t10k-labels-idx1-ubyte.gz
Label: [1]
Prediction: [1]
Label: [2]
Prediction: [2]
Label: [8]
Prediction: [8]
```



20

MNIST using MLP

■ This example shows how to use MLP to test MNIST

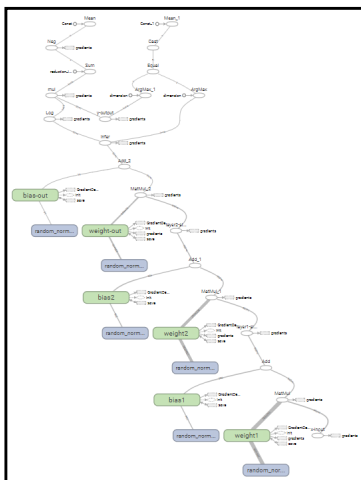
- ▶ Step 1: go to your project directory
 - ➔ [user@host] cd \$(PROJECT)/codes/tensorflow-project/mnist-project/mlp
- ▶ Step 2: see the codes
- ▶ Step 3: run Python under virtual environment
 - ➔ (do not forget to run '\$ source ~/tensorflow/bin/activate')
 - ➔ [user@host] python mnist_mlp.py
 - ➔ [user@host] python mnist_inference.py

```
[user@host] cd $(PROJECT)/codes/tensorflow-project/mnist-project/mlp
[user@host] python mnist_mlp.py
[user@host] python mnist_inference.py
```

21

MNIST using MLP

■ Make more structured style using "with tf.name.scope("LAYER1"):"




22

MNIST using single hidden layer network

■ Logistic regression for MNIST

- ▶ 784 (28x28) inputs of black and white → converted to floating number 0.0 ~ 1.0
- ▶ 10 outputs representing digit 0 to 9
- ▶ one hidden layer
- ▶ 256 features for each hidden layer



`$(PROJECT)/codes/tensorflow-projects/mnist-projects/single`

23

MNIST using single hidden layer network

■ Total testing patterns

- ▶ 60K images

■ Batch number

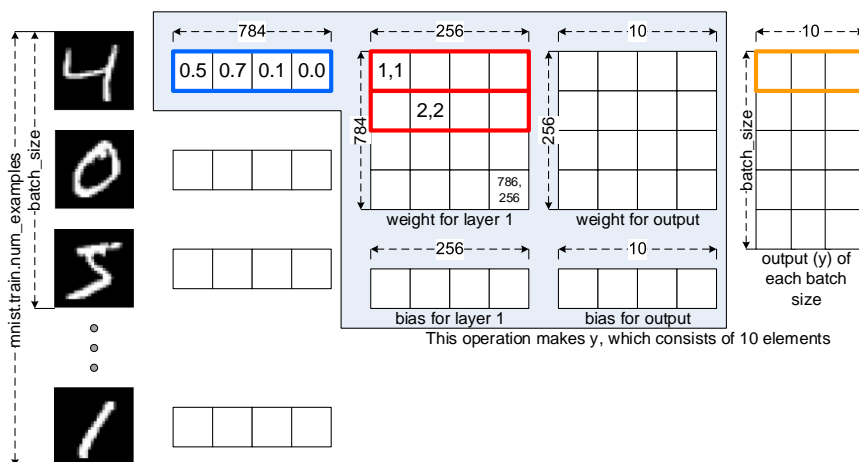
- ▶ 100 → group 600 batches

■ Epoch number

- ▶ The number of whole trainings (forward calculation and back propagation)

24

MNIST using single hidden layer network



25

MNIST using single hidden layer network

.... common parts are not shown

```
# tf Graph Input
x = tf.placeholder("float", [None, 784], name="x-input") # mnist data image of
shape 28*28
y = tf.placeholder("float", [None, 10], name="y-output") # 0-9 digits recognition =>
10 classes

# -----
# Model weights
W = tf.Variable(tf.zeros([784, 10]), name="weight")
b = tf.Variable(tf.zeros([10]), name="bias")

# -----
# inference -> hypothesis
infer = tf.nn.softmax(tf.matmul(x, W) + b, name="infer")

# Minimize error using cross entropy
cost = tf.reduce_mean(-tf.reduce_sum(y * tf.log(infer), reduction_indices=1))
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
# Evaluate -- 0~1 accuracy
correct = tf.equal(tf.argmax(infer, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct, "float"))

# Initializing the variables
init = tf.global_variables_initializer()
```

26

MNIST using single hidden layer network

```
#-----
with tf.Session() as sess:
    sess.run(init)
    # Training cycle
    for epoch in range(training_epochs):
        avg_cost = 0.; avg_accu = 0.
        total_batch = int( mnist.train.num_examples / batch_size )
        # Loop over all batches
        for i in range(total_batch):
            batch_xs, batch_ys = mnist.train.next_batch(batch_size)
            # Fit training using batch data
            sess.run(optimizer, feed_dict={x: batch_xs, y: batch_ys})
            # Compute average loss
            avg_cost += sess.run(cost, feed_dict={x: batch_xs, y: batch_ys})
            avg_accu += sess.run(accuracy, feed_dict={x: batch_xs, y: batch_ys})
        avg_cost /= total_batch; avg_accu /= total_batch
        # Display logs per epoch step
        if epoch % display_step == 0:
            print "Epoch:", "%04d" % (epoch + 1),\
                  "cost=", "{:.5f}".format(avg_cost),\
                  "accuracy=", "{:.5f}".format(avg_accu)
            if avg_accu > 0.95:
                break
        print "Optimization finished"
.... common parts are not shown ....
```

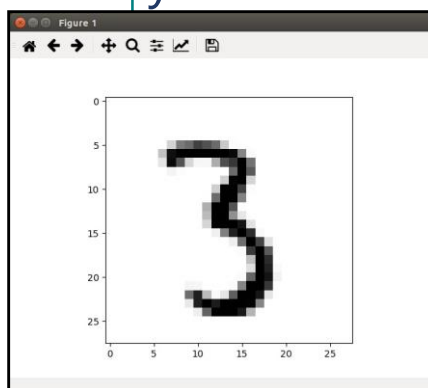
epoch times over all
data set

for each batch

27

MNIST using single hidden layer network

```
(tensorflow)$ python mnist_single.py
Extracting ../dataset/train-images-idx3-ubyte.gz
Extracting ../dataset/train-labels-idx1-ubyte.gz
Extracting ../dataset/t10k-images-idx3-ubyte.gz
Extracting ../dataset/t10k-labels-idx1-ubyte.gz
Epoch: 0001 cost= 10.12062 accuracy= 0.12273
Epoch: 0011 cost= 1.31484 accuracy= 0.74938
Epoch: 0021 cost= 0.97869 accuracy= 0.80429
Epoch: 0031 cost= 0.83487 accuracy= 0.82842
....
Epoch: 0071 cost= 0.61011 accuracy= 0.86305
Epoch: 0081 cost= 0.58191 accuracy= 0.86751
Epoch: 0091 cost= 0.55826 accuracy= 0.87131
....
Epoch: 0191 cost= 0.43436 accuracy= 0.89205
Epoch: 0201 cost= 0.42729 accuracy= 0.89313
Epoch: 0211 cost= 0.42070 accuracy= 0.89489
....
Epoch: 0471 cost= 0.33196 accuracy= 0.91184
Epoch: 0481 cost= 0.33010 accuracy= 0.91236
Epoch: 0491 cost= 0.32824 accuracy= 0.91265
Optimization finished
Accuracy: 0.9106
Label: [3]
Prediction: [3]
```



Training accuracy

Validation accuracy

Predict correct label

28

MNIST using single hidden layer network

■ This example shows how to use single hidden layer to test MNIST

- ▶ Step 1: go to your project directory
 - ➔ [user@host] cd \$(PROJECT)/codes/tensorflow-project/mnist-project/single
- ▶ Step 2: see the codes
- ▶ Step 3: run Python under virtual environment
 - ➔ (do not forget to run '\$ source ~/tensorflow/bin/activate')
 - ➔ [user@host] python mnist_single.py
 - ➔ [user@host] python mnist_inference.py

Prepare your own code in
\$(PROJECT)/codes/tensorflow-projects/mnist-projects/single


```
[user@host] cd $(PROJECT)/codes/tensorflow-project/mnist-project/single
[user@host] python mnist_single.py
[user@host] python mnist_inference.py
```

29

MNIST using CNN

■ Logistic regression for MNIST

- ▶ 784 (28x28) inputs of black and white → converted to floating number 0.0 ~ 1.0
- ▶ 10 outputs representing digit 0 to 9
- ▶ one hidden layer
- ▶ 256 features for each hidden layer




\$(PROJECT)/codes/tensorflow-projects/mnist-projects/cnn

30

MNIST using CNN: 1st convolutional layer

- Input: 28x28 pixels
- Convolution filter: 32 kernels with 5x5
- Convolution: stride 1
 - It generates 28x28x32 number of elements
- Result: 28x28



```
... ..
with tf.name_scope("WEIGHT-BIAS"):
    weights = {
        # 5x5 conv, 1 input, 32 outputs
        'wc1': tf.Variable(tf.random_normal([5, 5, 1, 32]),\
                               name="wc1"),
    }
    ... ..


with tf.name_scope("CONV-POOL-1"):
    # Convolution Layer
    conv1 = conv2d(x, weights['wc1'], biases['bc1'])
    ... ..

def conv2d(x, W, b, strides=1):
    with tf.name_scope("CONV"):
        # Conv2D wrapper, with bias and relu activation
        x = tf.nn.conv2d(x, W, strides=[1, strides, strides, 1],\
                           padding='SAME')
        x = tf.nn.bias_add(x, b)
    return tf.nn.relu(x)
```

31

MNIST using CNN: 1st pooling

- Input: 32 features with 28x28
- Max pooling filter: 5x5
- Convolution: stride 2
 - It generates 14x14x32 elements
- Results: 14x14




```
with tf.name_scope("CONV-POOL-1"):
    # Convolution Layer
    conv1 = conv2d(x, weights['wc1'], biases['bc1'])
    # Max Pooling (down-sampling)
    conv1 = maxpool2d(conv1, k=2)
```

32

MNIST using CNN: 2nd convolutional layer

- Input: 32 features with 14x14 pixels
- Convolution filter: 64 kernels with 5x5
- Convolution: stride 1
 - It generates 1/4 number of elements
- Results in 14x14x64




Conv: 5x5
64 filters
Padding: 2

```
... ..
with tf.name_scope("WEIGHT-BIAS"):
    weights = {
        # 5x5 conv, 1 input, 32 outputs
        'wc1': tf.Variable(tf.random_normal([5, 5, 1, 32])),
        name="wc1"),
        # 5x5 conv, 32 inputs, 64 outputs
        'wc2': tf.Variable(tf.random_normal([5, 5, 32, 64])),
        name="wc2"),
    ... ..
    with tf.name_scope("CONV-POOL-2"):
        # Convolution Layer
        conv2 = conv2d(conv1, weights['wc2'], biases['bc2'])
        # Max Pooling (down-sampling)
        conv2 = maxpool2d(conv2, k=2)
    ... ..
def conv2d(x, W, b, strides=1):
    with tf.name_scope("CONV"):
        # Conv2D wrapper, with bias and relu activation
        x = tf.nn.conv2d(x, W, strides=[1, strides, strides, 1],
            padding='SAME')
        x = tf.nn.bias_add(x, b)
    return tf.nn.relu(x)
```

33

MNIST using CNN: 2nd pooling

- Input: 64 features with 14x14
- Max pooling filter: 2x2
- Convolution: stride 2
 - It generates 1/4 number of elements
- Results in 7x7x64




Pool: 2x2
stride: 2

```
with tf.name_scope("CONV-POOL-2"):
    # Convolution Layer
    conv2 = conv2d(conv1, weights['wc2'], biases['bc2'])
    # Max Pooling (down-sampling)
    conv2 = maxpool2d(conv2, k=2)
```

34

MNIST using CNN: fully connected layer

- Input: 64 features with 7x7
- Reshaping: 3-D array to 1-D vector
 - ▶ $64 \times 7 \times 7 \rightarrow 3136$
- Neurons:




```
with tf.name_scope("FC"):
    # Fully connected layer
    # Reshape conv2 output to fit fully connected layer
    input
    fc1 = tf.reshape(conv2, [-1,
weights['wd1'].get_shape().as_list()[0]])
    fc1 = tf.add(tf.matmul(fc1, weights['wd1']),
    biases['bd1'])
    fc1 = tf.nn.relu(fc1)
    # Apply Dropout
    fc1 = tf.nn.dropout(fc1, dropout)
```

35

MNIST using CNN: read-out layer


- Input: 1024 neurons
- Output: 10 classes



```
with tf.name_scope("OUT"):
    # Output, class prediction
    out = tf.add(tf.matmul(fc1, weights['out']), biases['out'],
name="add-output")
```

36

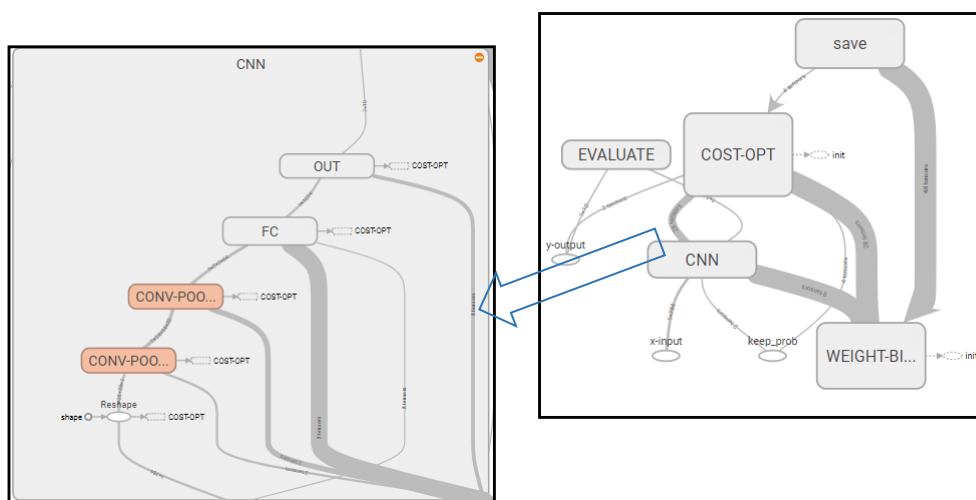
MNIST using CNN



<http://www.cnblogs.com/BigBallon/p/6701846.html>

37

MNIST using CNN



38

MNIST using CNN

■ This example shows how to use CNN to test MNIST

- ▶ Step 1: go to your project directory
 - ➔ [user@host] `cd $(PROJECT)/codes/tensorflow-project/mnist-project/cnn`
- ▶ Step 2: see the codes
- ▶ Step 3: run Python under virtual environment
 - ➔ (do not forget to run '\$ source ~/tensorflow/bin/activate')
 - ➔ [user@host] `python mnist_cnn.py`
 - ➔ [user@host] `python mnist_inference.py`

```
[user@host] cd $(PROJECT)/codes/tensorflow-project/mnist-project/cnn
[user@host] python mnist_cnn.py
[user@host] python mnist_inference.py
```

39

(주)퓨처디자인시스템

34051 대전광역시 유성구 문지로 193, KAIST 문지캠퍼스, F723호
(042) 864-0211~0212 / contact@future-ds.com / www.future-ds.com

Future Design Systems, Inc.

Faculty Wing F723, KAIST Munji Campus, 193 Munji-ro, Yuseong-gu, Daejeon 34051, Korea
+82-042-864-0211~0212 / contact@future-ds.com / www.future-ds.com



FUTURE
Design Systems